

NIC
EMPOWER

November 13-15, Oslo Spektrum

Karim El-Melhaoui, O3 Cyber

Azure Security Assessments: Lessons from the Field

Quick intro

- Karim El-Melhaoui (karim@o3c.no)
- Principal at O3 Cyber (o3c.no)
- CloudSec Researcher
- Azure MVP but I love **most** clouds

This is a deep dive session

1. Introduction to assessment and Scoping
2. Discovery, Data Collection and Tooling
3. Writing our own Tools
4. Attack Paths
5. Research Release: Abusing Azure Data Factory
6. Research Preview: OIDC Attack Path Analyzer

Introduction to a Cloud Security Assessment

What is an assessment

- Infrastructure and Architecture review
 - Access and Authentication assessment
 - Configuration Analysis
 - Exposure and Vulnerabilities

Goal

- Uncover critical design flaws
- Identify Toxic Combinations (Attack Paths)
- Assess the efficacy of controls
- Find exposed vulnerabilities

Scoping an assessment: Technical

- Service Coverage
 - Azure Resource Manager
 - Microsoft Graph
 - Azure Services
 - GitHub / Azure DevOps / CI tool
- Platform usage

Scoping an assessment: Organizational

- Business Context
- Expectations

What we need to know

- Third-parties we should expect to see
- Privileged Access procedures
- Segmentation policy
- Guardrails / Governance procedures
- Tooling

Permissions

- Non-intrusive:
 - Global Reader
 - Reader on Tenant Root Group
 - Reader in GitHub
- Intrusive: Same as a 'DevOps' user

Discovery / Data Collection

Open-source tools

Tools

PROWLER




Prowler



- Multi-Cloud
- Built in Python with SDK's
- Compliance Benchmarks ++

Compliance scan with Prowler



(nicconf_demo) karimel-melhaoui@mbp nicconf %

Why do we complicate things beyond this?

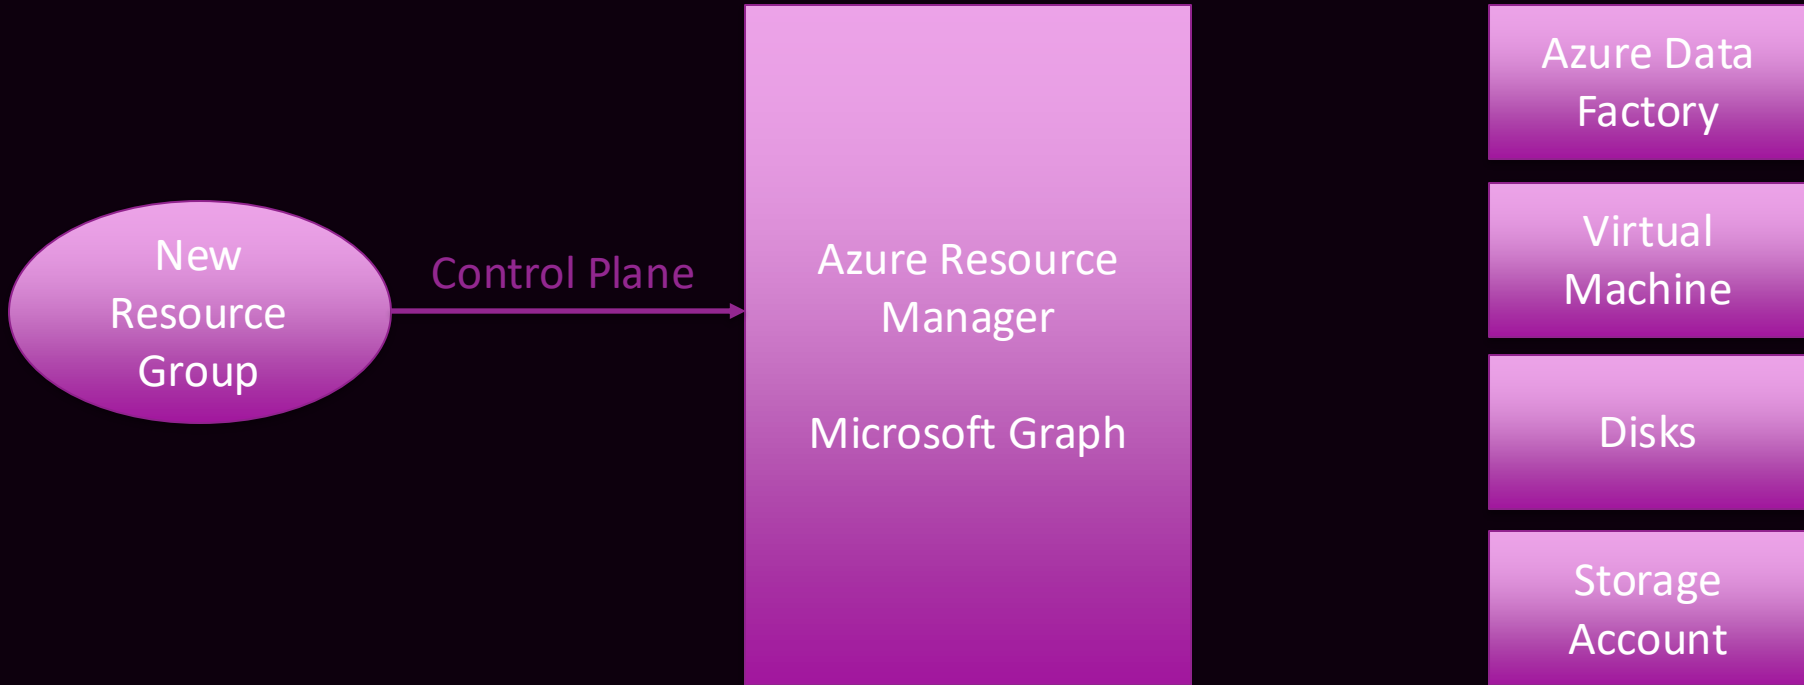
What the client **shouldn't** care about

- Compliance for the sake of being compliant
- Unencrypted disks
- Isolated configuration findings

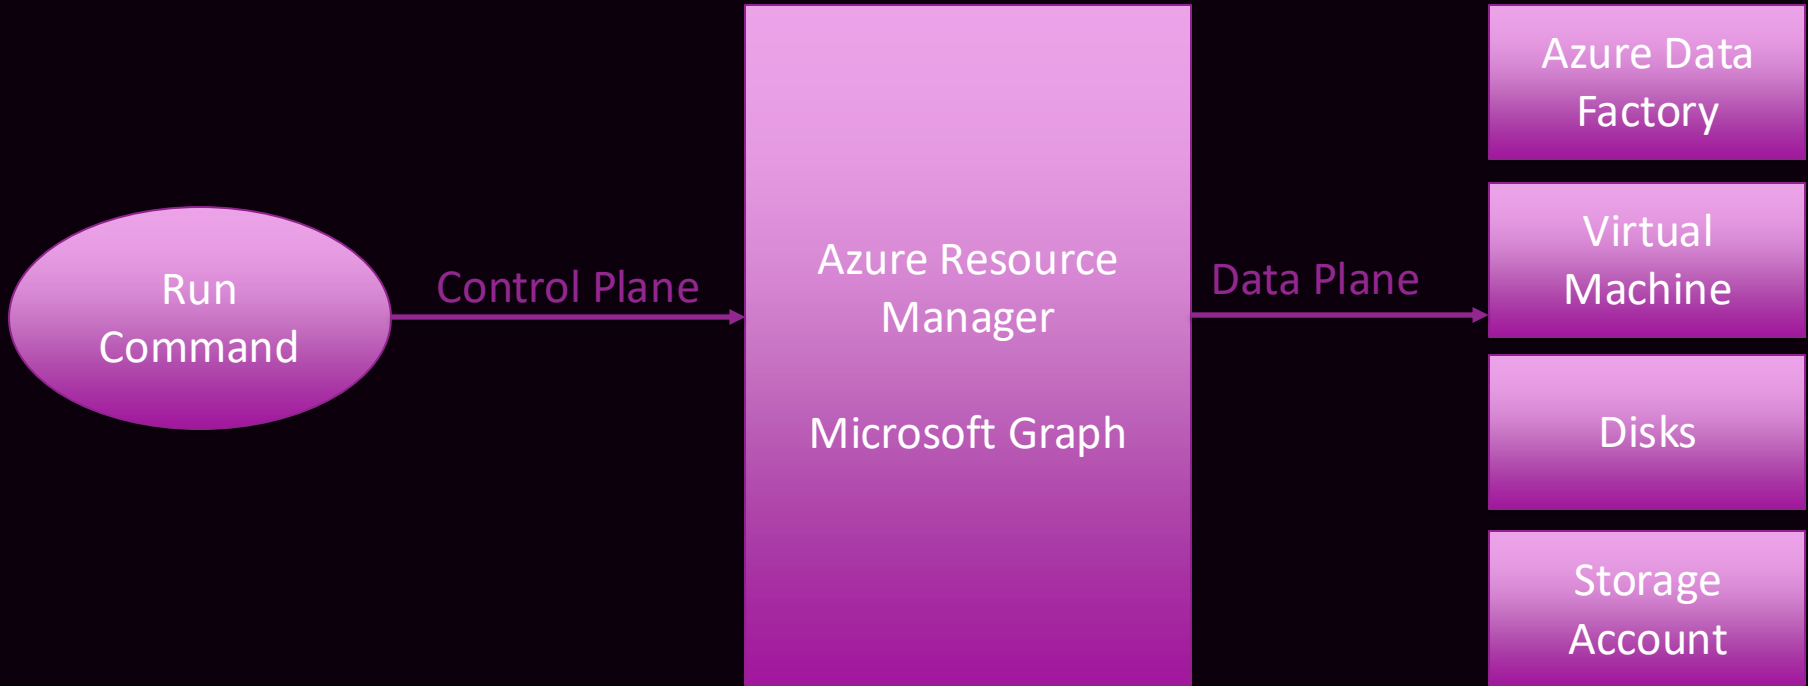
* For the record, I love Prowler and use it all the time!

Writing our own tools

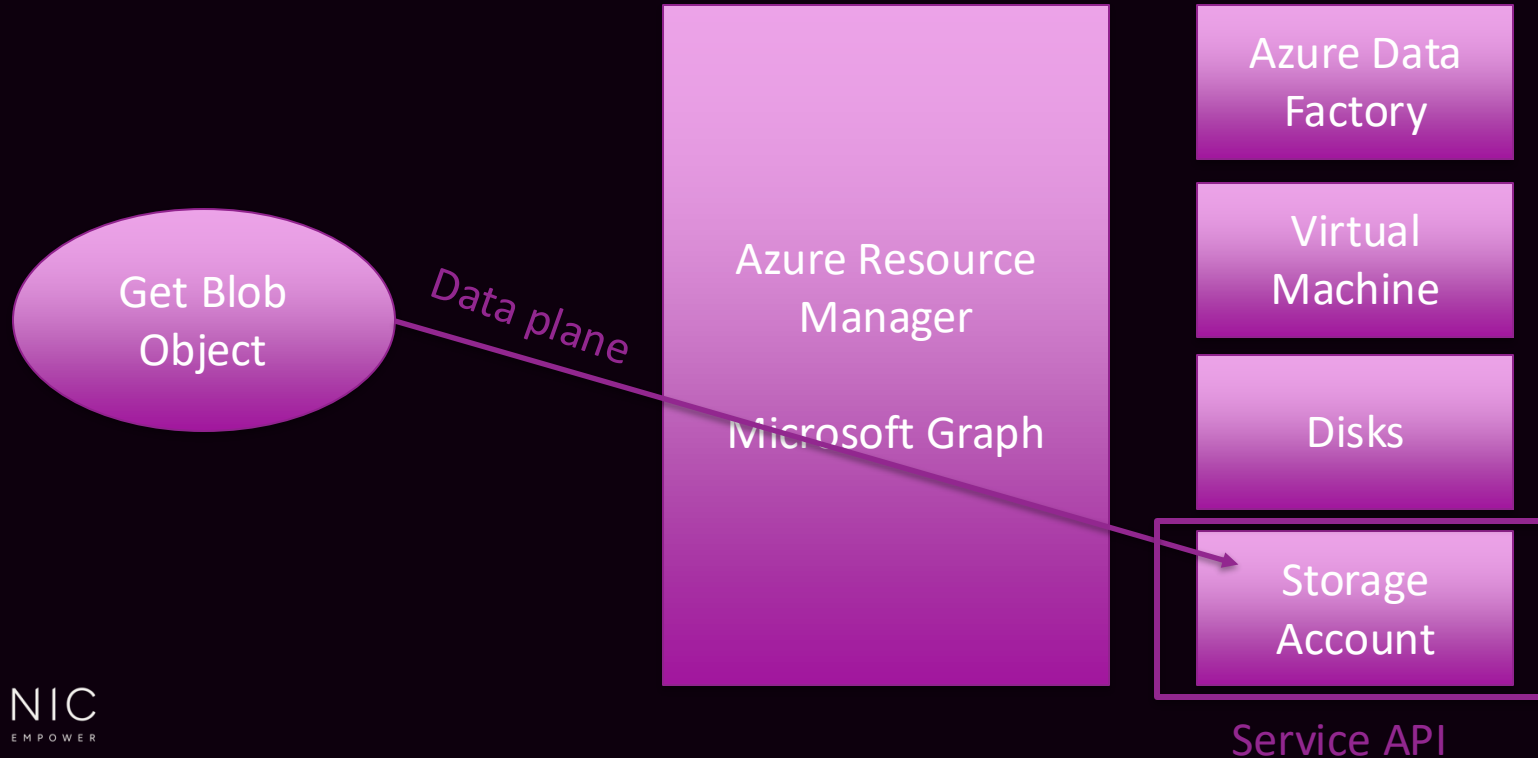
Control Plane vs. Data Plane



Control Plane vs. Data Plane



Control Plane vs. Data Plane



REST

Tag
Resource
Group



```
graph LR; A([Tag Resource Group]) --> B[Header<br/>Operation<br/>URI<br/>Request]; B --> C[Azure Resource Manager<br/>Microsoft Graph];
```

The diagram illustrates a REST API request flow. It begins with an oval labeled 'Tag Resource Group'. An arrow points from this oval to a rounded rectangle containing the components of a REST request: 'Header', 'Operation', 'URI', and 'Request'. A second arrow points from this rectangle to a large rectangle on the right, which lists the services: 'Azure Resource Manager' and 'Microsoft Graph'.

Header

Operation

URI

Request

Azure Resource
Manager
Microsoft Graph

REST

- Full flexibility and control
- Language independent
- Low-Level operations

SDK

Tag
Resource
Group

Package

Client

Method

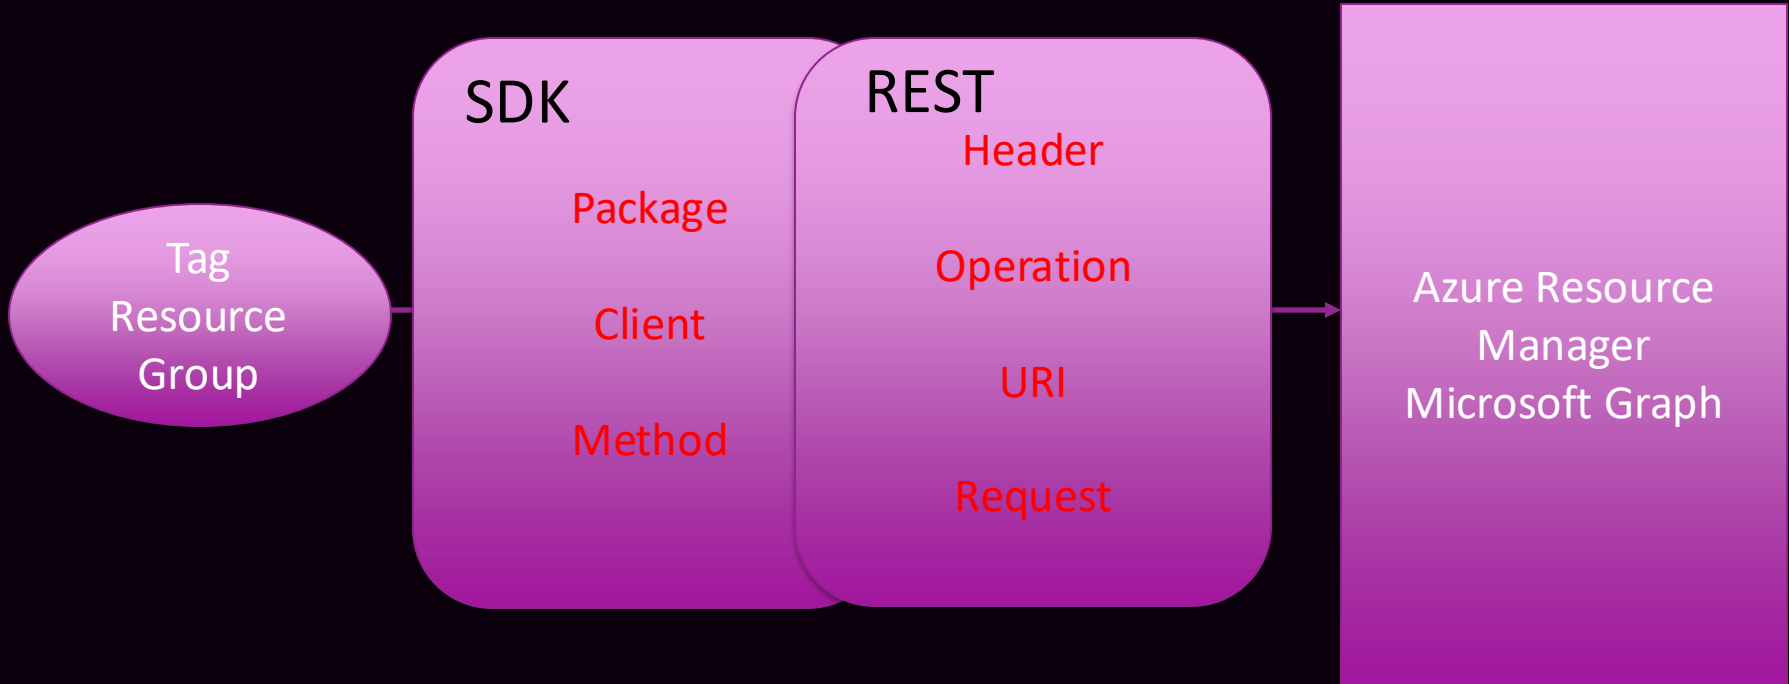
REST

Azure Resource
Manager
Microsoft Graph

SDK

- Built-in Authentication Handling
- Automatic retry and error handling
- Language specific

Best of Both Worlds



End Goal



Authentication

import msal

```
class AuthClientBase:

    def __init__(self, client_id, client_credential, tenant_id, scope):
        self.app = msal.ConfidentialClientApplication(
            client_id=client_id,
            client_credential=client_credential,
            authority=f"https://login.microsoftonline.com/{tenant_id}",
        )
        self.scope = scope

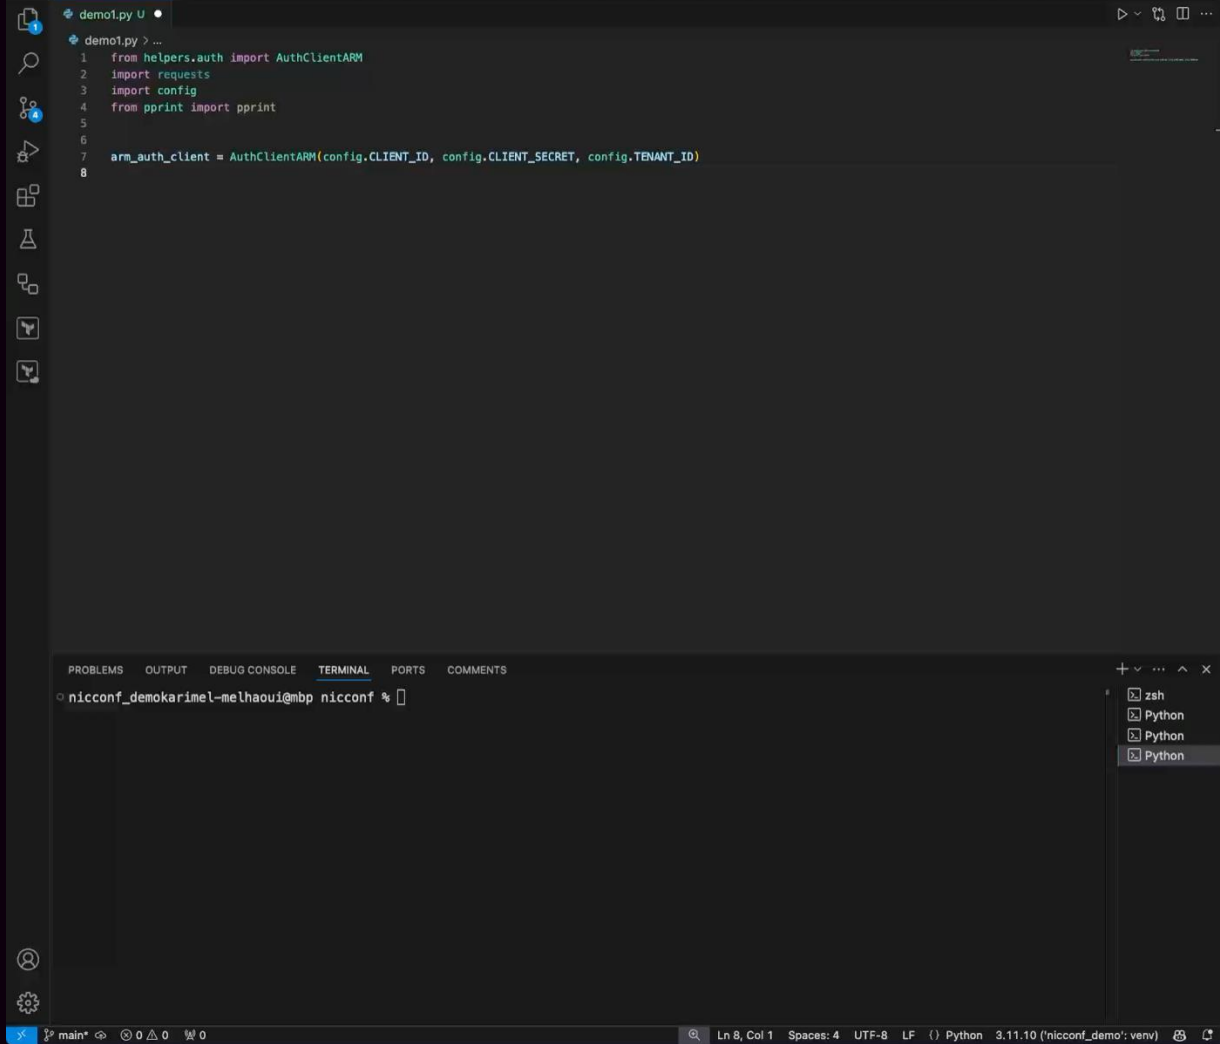
    def get_token(self):
        result = self.app.acquire_token_silent(scopes=[self.scope], account=None)
        if not result:
            result = self.app.acquire_token_for_client(scopes=[self.scope])

        if "access_token" in result:
            return result["access_token"]
        else:
            raise Exception("Failed to obtain access token")
```


Authentication

```
class AuthClientGraph(AuthClientBase):  
    def __init__(self, client_id, client_credential, tenant_id):  
        super().__init__(  
            client_id,  
            client_credential,  
            tenant_id,  
            "https://graph.microsoft.com/.default",  
        )
```

```
class AuthClientARM(AuthClientBase):  
    def __init__(self, client_id, client_credential, tenant_id):  
        super().__init__(  
            client_id,  
            client_credential,  
            tenant_id,  
            "https://management.azure.com/.default",  
        )
```



Resource Groups

```
def get_resource_groups(arm_auth_client, subscription):  
    """  
    Fetch the list of resource groups for a specific subscription from the Azure Management API.  
  
    Args:  
        arm_auth_client: The authentication client to use for fetching the token.  
        subscription (str): The subscription ID to fetch resource groups for.  
  
    Returns:  
        List[Dict]: A list of dictionaries containing the resource group details.  
    """  
    token = arm_auth_client.get_token()  
    url = f"https://management.azure.com/{subscription}/resourceGroups?api-version=2020-01-01"  
    headers = {"Authorization": f"Bearer {token}"}  
    response = requests.get(url, headers=headers)  
    if response.status_code == 200:  
        return response.json().get("value", [])  
    else:  
        print(  
            f"Error fetching resource groups: {response.status_code} - {response.text}"  
        )  
        return []
```

Resource Group Role Assignments

```
def get_rg_role_assignment(arm_auth_client, subscription, resource_group):  
    """  
    Fetch the list of role assignments for a specific resource group from the Azure Management API.  
  
    Args:  
        arm_auth_client: The authentication client to use for fetching the token.  
        subscription (str): The subscription ID to fetch role assignments for.  
        resource_group (str): The resource group ID to fetch role assignments for.  
  
    Returns:  
        List[Dict]: A list of dictionaries containing the role assignment details.  
    """  
    token = arm_auth_client.get_token()  
    url = f"https://management.azure.com/{resource_group}/providers/Microsoft.Authorization/roleAssignments?api-version=2022-04-01"  
    headers = {"Authorization": f"Bearer {token}"}  
    response = requests.get(url, headers=headers)  
    if response.status_code == 200:  
        return response.json().get("value", [])  
    else:  
        print(  
            f"Error fetching resource group role assignments: {response.status_code} - {response.text}"  
        )  
        return []
```

Resources

```
def get_resources(arm_auth_client, subscription, resource_group):  
    """  
    Fetch the list of resources for a specific resource group from the Azure Management API.  
  
    Args:  
        arm_auth_client: The authentication client to use for fetching the token.  
        subscription (str): The subscription ID to fetch resources for.  
        resource_group (str): The resource group ID to fetch resources for.  
  
    Returns:  
        List[Dict]: A list of dictionaries containing the resource details.  
    """  
    token = arm_auth_client.get_token()  
    url = f"https://management.azure.com/{resource_group}/resources?api-version=2021-04-01"  
    headers = {"Authorization": f"Bearer {token}"}  
    response = requests.get(url, headers=headers)  
    if response.status_code == 200:  
        return response.json().get("value", [])  
    else:  
        print(f"Error fetching resources: {response.status_code} - {response.text}")  
        return []
```

What if we want to get ALL DATA?

Live demo: Getting all the data

demo1.py U

{} role_assignments_processed.json X

output > {} role_assignments_processed.json > {} 0 > {} properties

```
1 {
2   {
3     "properties": {
4       "roleDefinitionId": "/subscriptions/b837386d-5ba9-4ff9-b55e-7b483e17b8bd/providers/Microsoft.Authorization/roleDefinitions/b64e21ea-ac4e-4cdf-9dc9-5b892992bee7",
5       "principalId": "709e8357-f0e2-4a9e-af0e-3ef2dd37378b",
6       "principalType": "ServicePrincipal",
7       "scope": "/subscriptions/b837386d-5ba9-4ff9-b55e-7b483e17b8bd",
8       "condition": null,
9       "conditionVersion": null,
10      "createdOn": "2021-01-02T03:38:50.3147500Z",
11      "updatedOn": "2021-01-02T03:38:50.3147500Z",
12      "createdBy": "d84f77b0-4e47-4a89-ad5b-1733c08e6752",
13      "updatedBy": "d84f77b0-4e47-4a89-ad5b-1733c08e6752",
14      "delegatedManagedIdentityResourceId": null,
15      "description": null,
16      "roleName": "Azure Connected Machine Onboarding"
17    },
18    "id": "/subscriptions/b837386d-5ba9-4ff9-b55e-7b483e17b8bd/providers/Microsoft.Authorization/roleAssignments/95b7fd60-da85-4a85-b6a4-5efa3434ee71",
19    "type": "Microsoft.Authorization/roleAssignments",
20    "name": "95b7fd60-da85-4a85-b6a4-5efa3434ee71"
21  },
22  {
23    "properties": {
24      "roleDefinitionId": "/subscriptions/b837386d-5ba9-4ff9-b55e-7b483e17b8bd/providers/Microsoft.Authorization/roleDefinitions/8e3af657-a8ff-443c-a75c-2fe8c4bcb635",
25      "principalId": "d84f77b0-4e47-4a89-ad5b-1733c08e6752",
26      "principalType": "User",
27      "scope": "/subscriptions/b837386d-5ba9-4ff9-b55e-7b483e17b8bd",
28      "condition": null,
29      "conditionVersion": null,
30      "createdOn": "2022-09-29T15:26:12.9136497Z",
31      "updatedOn": "2022-09-29T15:26:12.9136497Z",
32      "createdBy": "d84f77b0-4e47-4a89-ad5b-1733c08e6752",
33      "updatedBy": "d84f77b0-4e47-4a89-ad5b-1733c08e6752",
34      "delegatedManagedIdentityResourceId": null,
35      "description": null,
36      "roleName": "Owner"
37    },
38    "id": "/subscriptions/b837386d-5ba9-4ff9-b55e-7b483e17b8bd/providers/Microsoft.Authorization/roleAssignments/d2bff257-7c22-4e3e-a595-0a06cdad33",
```

main*

Ln 9, Col 32 Spaces: 2 UTF-8 LF JSON

I could have just used PowerShell or Az cli?

... No. Because

- Full control of what I want to request
- Returns the full object
- Extensible
- Interact outside of what Microsoft wants you to.


Classic Admins in SDK? Nope.

```
def get_classic_admins(arm_auth_client, subscription):
    token = arm_auth_client.get_token()
    url = f"https://management.azure.com{subscription}/providers/Microsoft.Authorization/classicAdministrators?api-version=2015-07-01"
    headers = {"Authorization": f"Bearer {token}"}
    response = requests.get(url, headers=headers)
    if response.status_code == 200:
        return response.json().get("value", [])
    else:
        print(
            f"Error fetching subscription role assignments: {response.status_code} - {response.text}"
        )
        return []
```

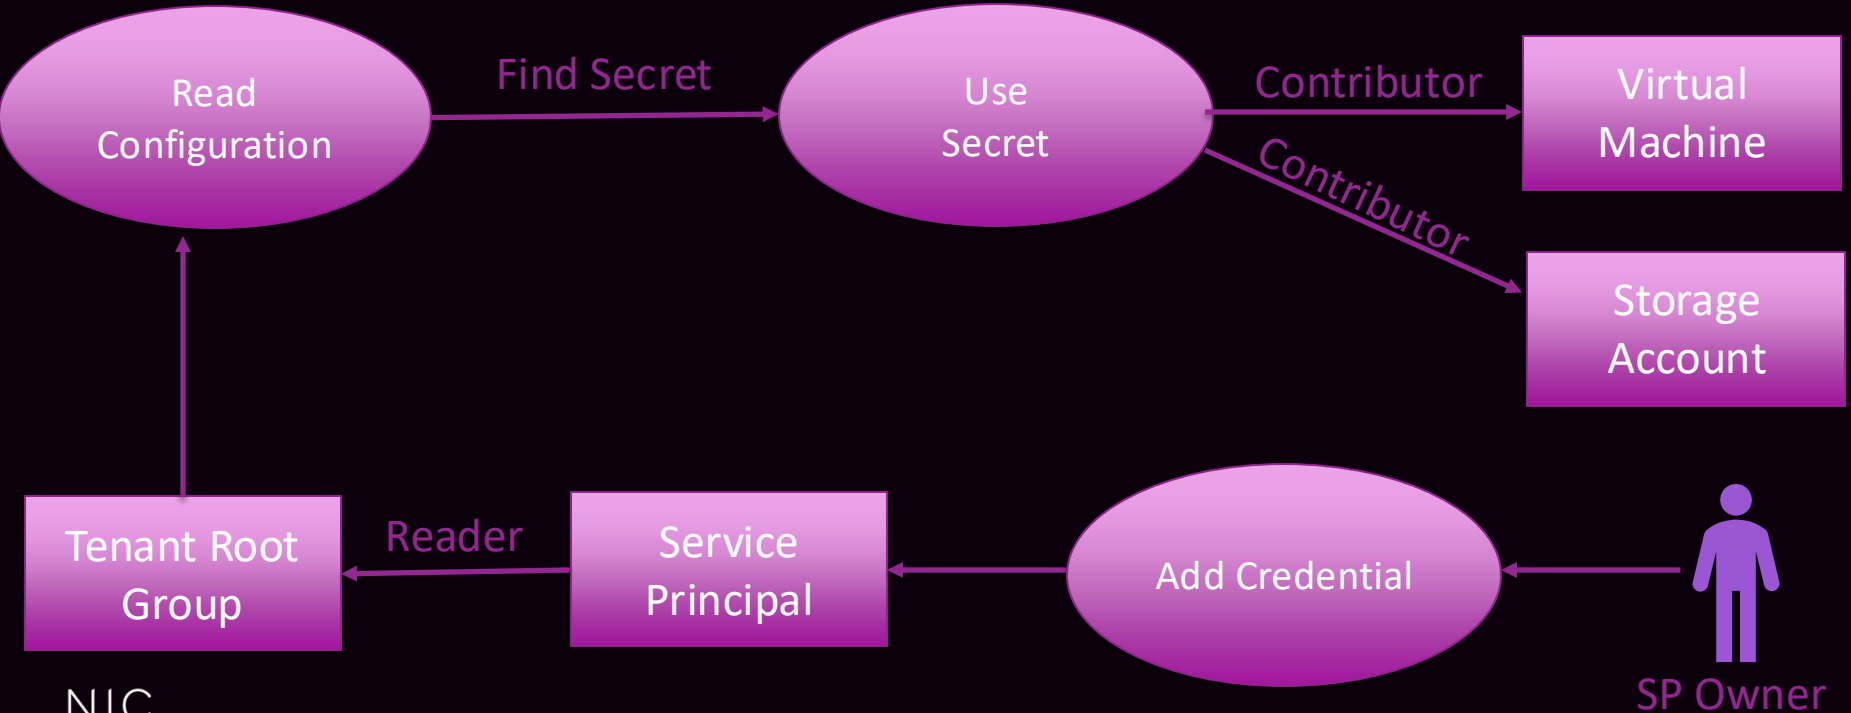
Let's dive deeper.

Attack Path Introduction

Simple Attack Path




Attack Path



What causes an Attack Path?

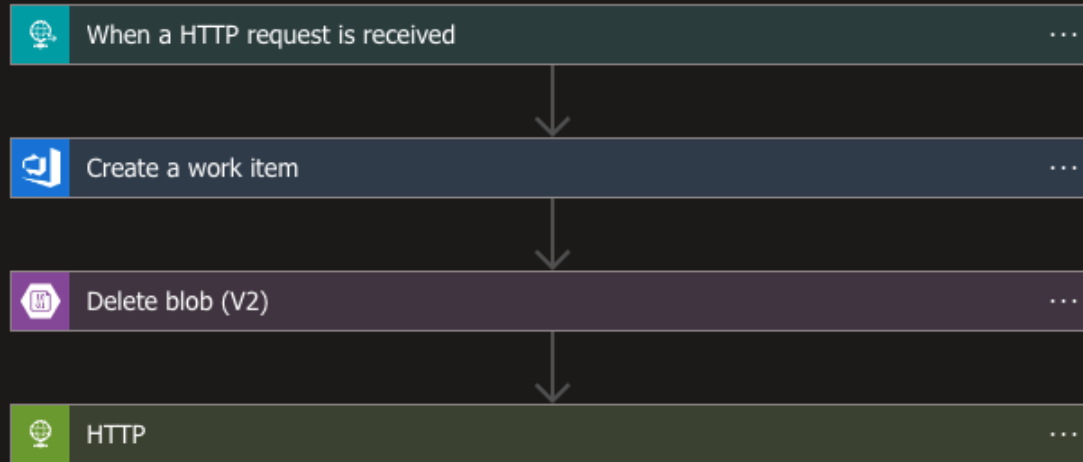
- Serverless with privileged Managed Identities
- Clear-text secrets
- Roles such as Application Administrator, Helpdesk Admin
- Users that are Owners on applications
- Exposed endpoints

Attack Path Tools



Finding Secrets

Logic Apps



+ New step

Logic Apps

```
def get_logic_apps_configuration(arm_auth_client, subscription, resource_group):  
    """  
    Fetch the configuration of Logic Apps for a specific resource group from the Azure Management API.  
  
    Args:  
        arm_auth_client: The authentication client to use for fetching the token.  
        subscription (str): The subscription ID to fetch Logic Apps for.  
        resource_group (str): The resource group ID to fetch Logic Apps for.  
  
    Returns:  
        List[Dict]: A list of dictionaries containing the Logic Apps configuration details.  
    """  
    token = arm_auth_client.get_token()  
    url = f"https://management.azure.com/{resource_group}/providers/Microsoft.Logic/workflows?api-version=2019-05-01"  
    headers = {"Authorization": f"Bearer {token}"}  
    response = requests.get(url, headers=headers)  
    if response.status_code == 200:  
        return response.json().get("value", [])  
    else:  
        print(f"Error fetching Logic Apps configuration: {response.status_code} - {response.text}")  
        return []
```

← →

nicconf

🔍

arm_data.py M

main.py M

main.py > ...

...

182 def main():

183 """

184 Main function to orchestrate the fetching and processing of data.

185 """

186 try:

187 graph_auth_client = AuthClientGraph(

188 config.CLIENT_ID, config.CLIENT_SECRET, config.TENANT_ID

189)

190 arm_auth_client = AuthClientARM(

191 config.CLIENT_ID, config.CLIENT_SECRET, config.TENANT_ID

192)

193

194 logger.info("Fetching subscriptions...")

195 subs = fetch_subscriptions(arm_auth_client)

196 pprint(subs)

197 with open("output/subscriptions.json", "w") as f:

198 json.dump(subs, f)

199

200

201 logger.info("Fetching Logic Apps configuration...")

202 logic_apps = fetch_logic_apps(arm_auth_client, subs)

203 pprint(logic_apps)

204 with open("output/logic_apps.json", "w") as f:

205 json.dump(logic_apps, f)

206

207 except Exception as e:

208 logger.error(f"An unexpected error occurred: {str(e)}")

209

210

211 if __name__ == "__main__":

212 main()

> password: As

PROBLEMS

OUTPUT

DEBUG CONSOLE

TERMINAL

PORTS

COMMENTS

nicconf_demokari:mel-melhaoui@mbp nicconf %

Ln 209, Col 1

Spaces: 4

UTF-8

This is ALWAYS the case.

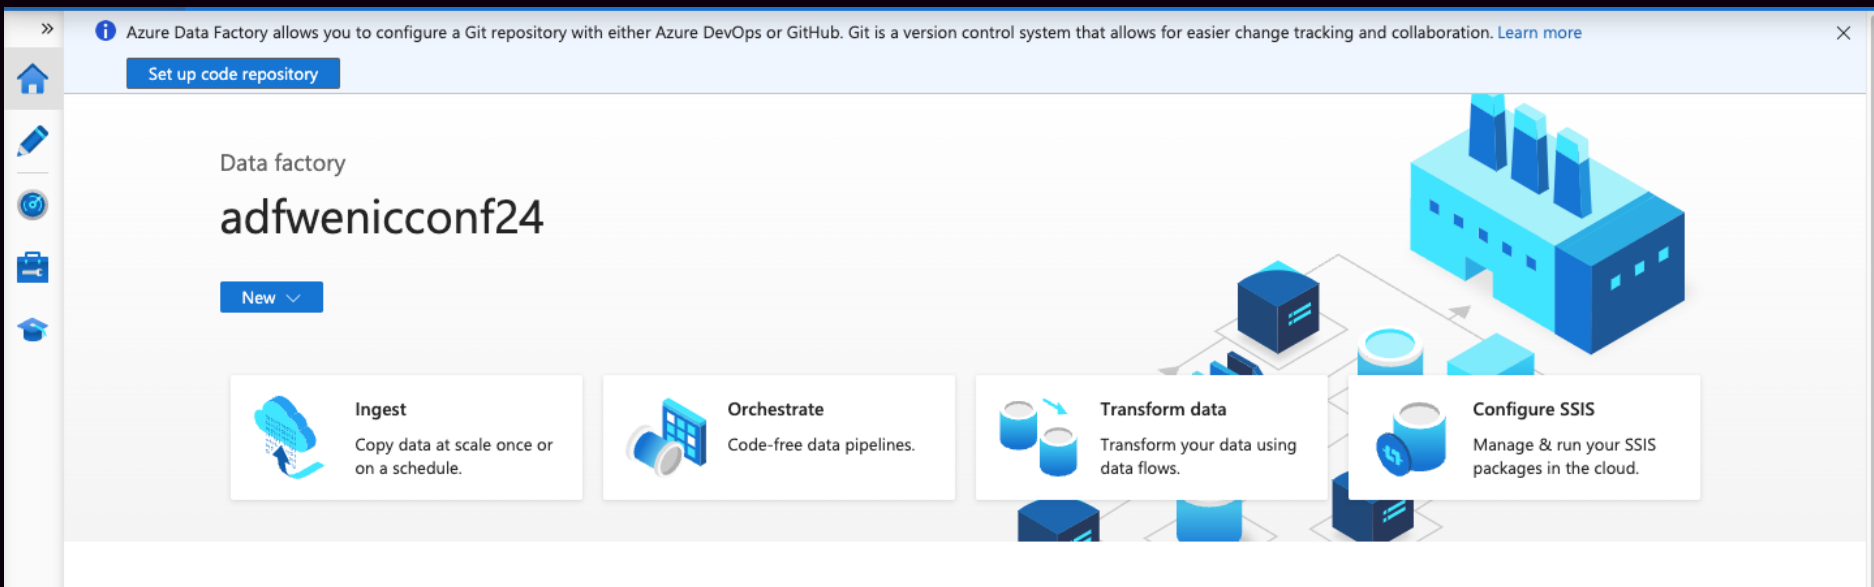
More Secrets

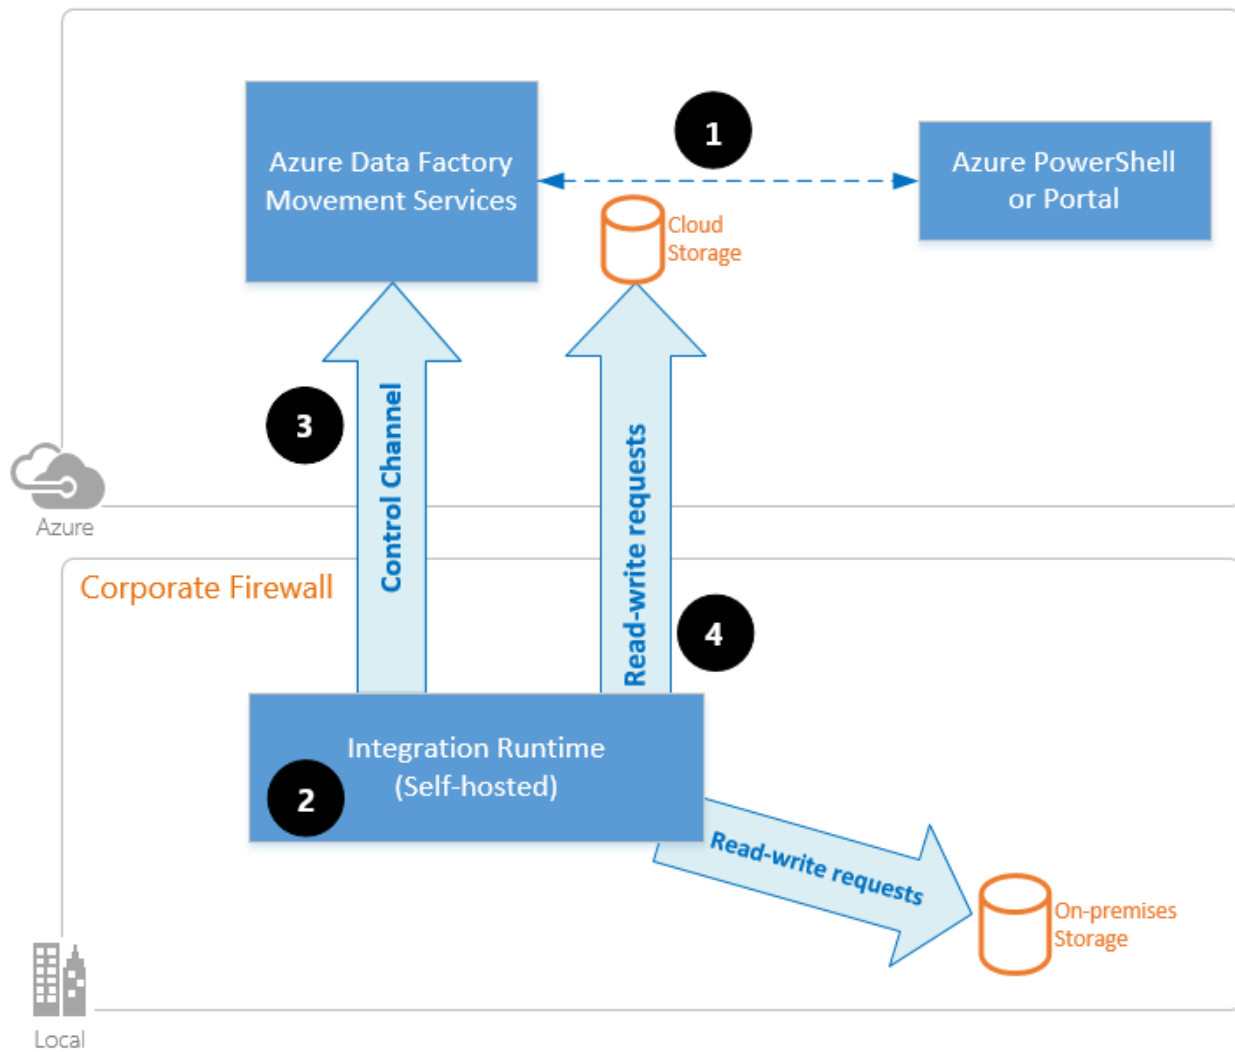


More secrets, means more attack paths

Abusing Azure Data Factory

Azure Data Factory





Azure Data Factory

- Azure Data Factory can execute jobs through Self-Hosted Integration Runtime
- Self-Hosted Integration Runtime Downloads Secrets

Exploiting SHIR

SHIR

The screenshot displays the Microsoft Integration Runtime Configuration Manager window. The title bar reads "Microsoft Integration Runtime Configuration Manager". The navigation menu includes "Home", "Settings", "Diagnostics", "Update", and "Help". The "Home" tab is active, showing a green checkmark icon and the message "Self-hosted node is connected to the cloud service". Below this, the configuration details are listed: "Data Factory: adfwenicconf24", "Integration Runtime: integrationRuntime1", and "Node: nicconf-shir". A "Stop Service" button is present. The "Data Source Credential" section, marked with an information icon, shows "Credential store: On-premises", "Credential status: In sync", and "Last backup time: N/A". It includes "Generate Backup" and "Import Backup" buttons. The status bar at the bottom indicates "Connected to the cloud service (Data Factory V2)" with a green checkmark and a refresh icon.

Microsoft Integration Runtime Configuration Manager

Home Settings Diagnostics Update Help

✓ Self-hosted node is connected to the cloud service

Data Factory: adfwenicconf24
Integration Runtime: integrationRuntime1
Node: nicconf-shir

Stop Service

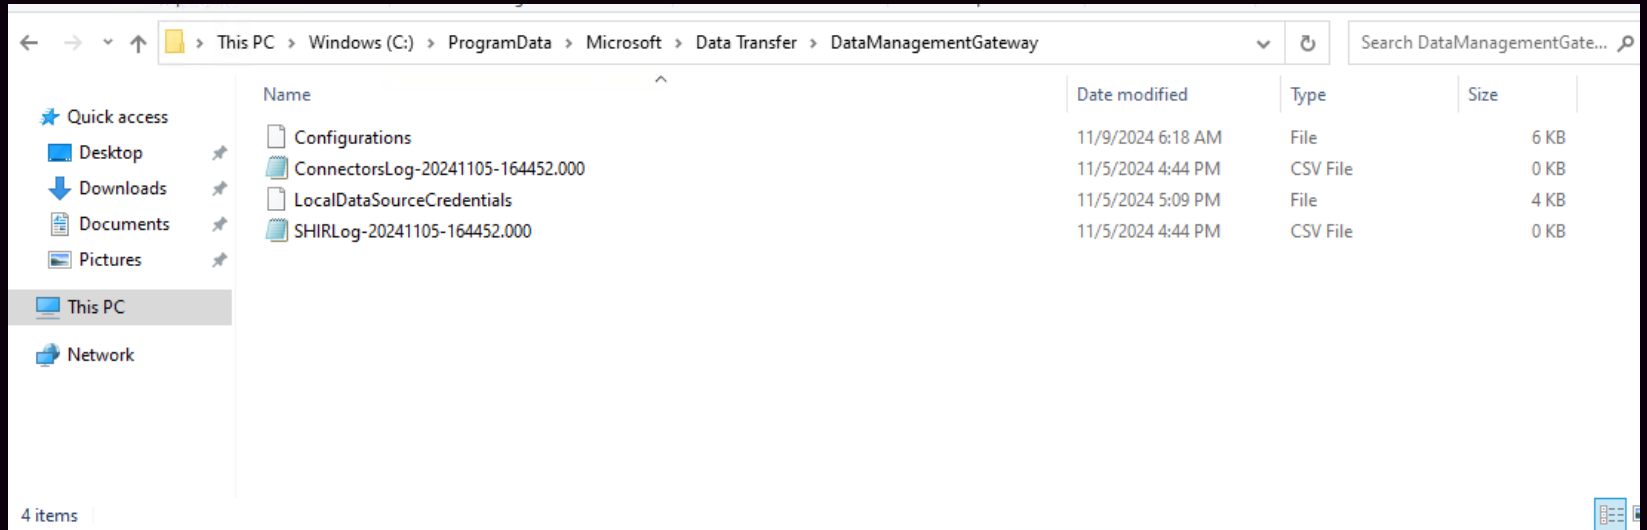
Data Source Credential ⓘ

Credential store: On-premises
Credential status: In sync
Last backup time: N/A

Generate Backup Import Backup

✓ Connected to the cloud service (Data Factory V2)

SHIR





Recycle Bin



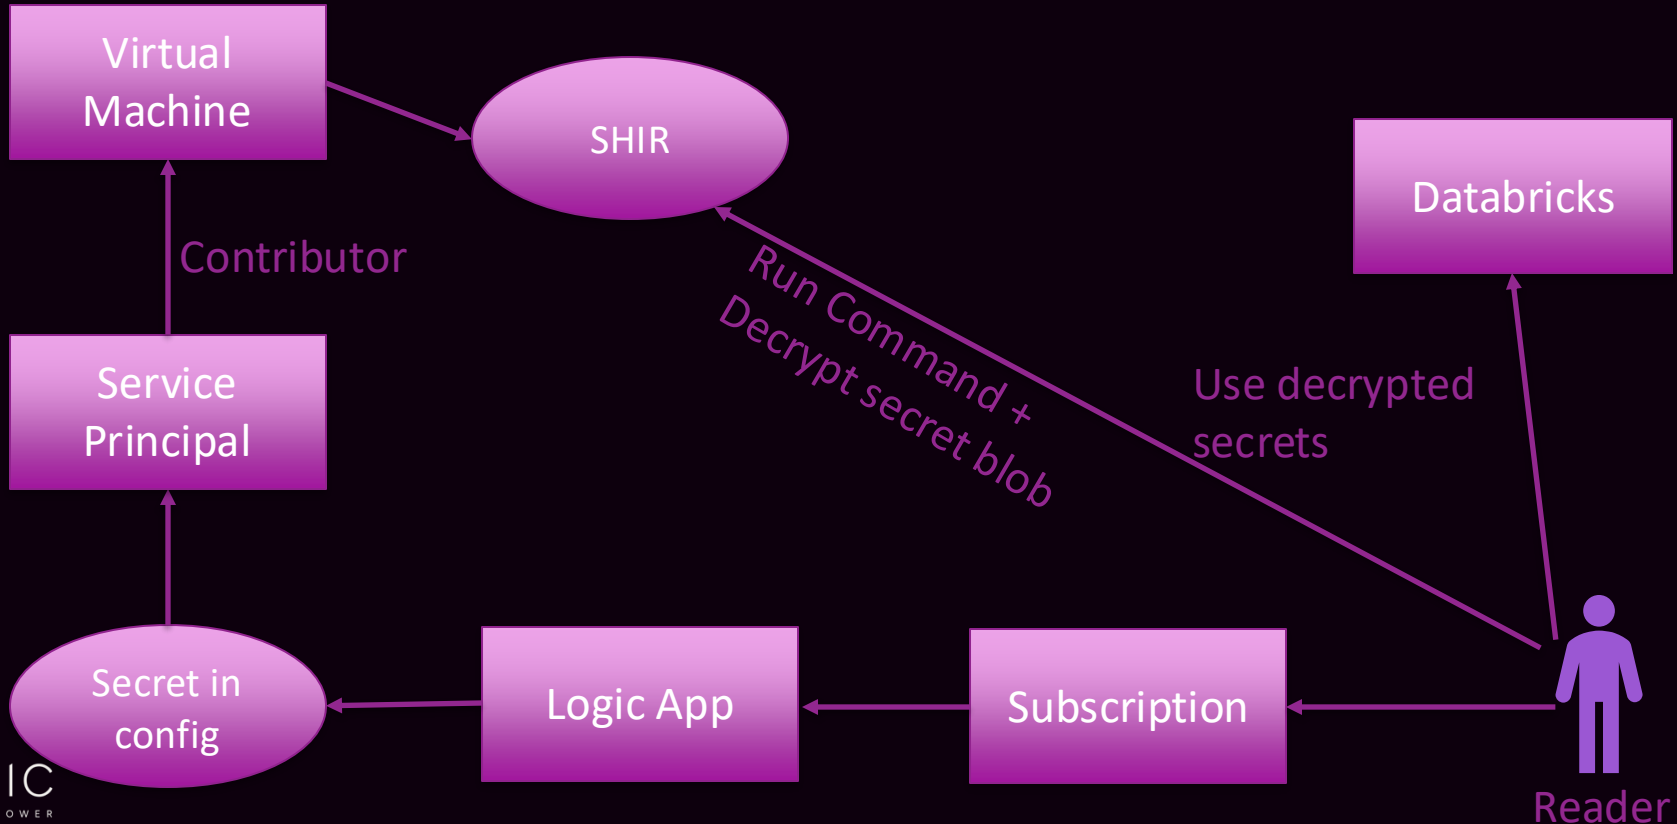
Microsoft
Edge



dddd



Data Factory



Azure Data Factory - Abusing the Self-Hosted Integration Runtime


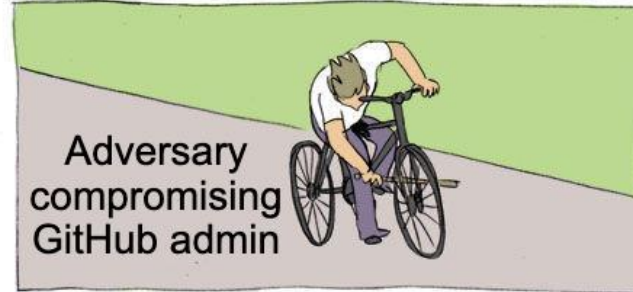

During a recent Cloud Security Assessment, we obtained **Contributor** permissions on a Virtual Machine (VM) configured as a Self-Hosted Integration Runtime (SHIR) instance in an Azure Subscription. This setup raised a question: could we leverage our access to pivot toward more valuable data? To explore this, we recreated the SHIR configuration in our lab environment and used Azure's **Run Command** to modify settings and gain full access to the host machine.

Observing SHIR Credentials with Procmon

Once inside the SHIR host, we ran **Procmon** to monitor activity as we created a new connection in Azure Data Factory that used the SHIR instance. During this process, we observed the creation of a file named `LocalDataSourceCredentials` in the following location:

```
C:\Windows\ProgramData\Microsoft\Data Transfer\DataManagementGateway
```

Going from Code to Cloud



Open ID Connect



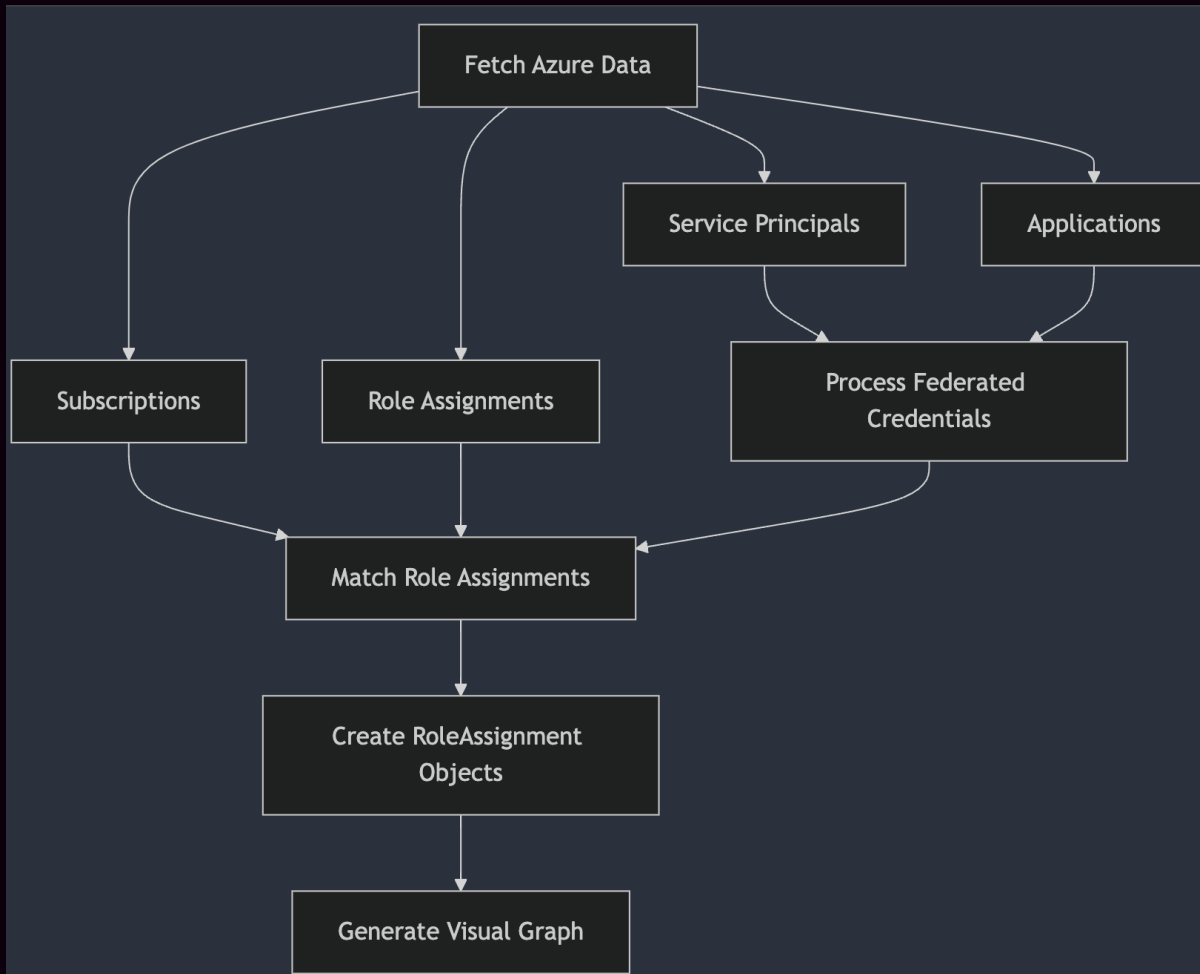
OIDC Config flaw

Issuer ⓘ	<input type="text" value="https://token.actions.githubusercontent.com"/> Edit (optional)
Organization *	<input type="text" value="karimelmel"/>
Repository *	<input type="text" value="cloud-infra-as-code"/>
Entity type	<input type="text" value="Branch"/> ▼
Based on selection *	<input type="text" value="main"/>
Subject identifier * ⓘ	<input type="text" value="repo:karimelmel/cloud-infra-as-code:ref:refs/heads/main"/> Generate this value using your GitHub account details instead

repo:org/repo:pull_request

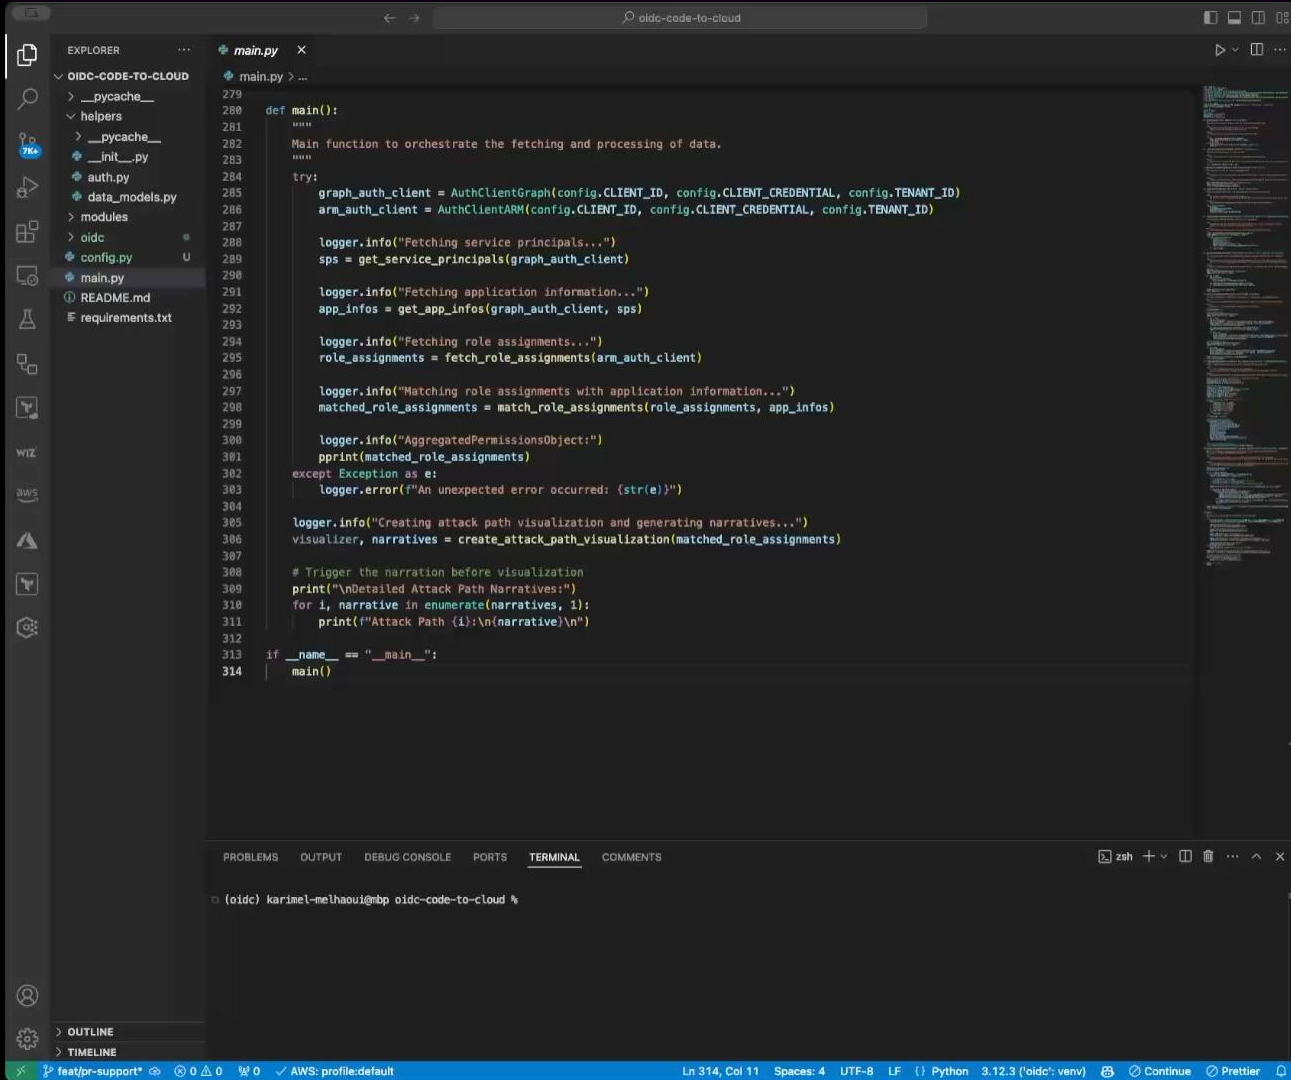
repo:org/*

repo:org/repo:<branch>




Finding all your misconfigured OIDC

*Full session and tool release scheduled for Q1 '25



Attack Path: GitHub to Azure via Entra ID



Reporting

Report output

- Impact
- Reproducibility
- Technical details vs. executive summary
- Actionable improvements, not one-click remediations

Thank you for watching!

karim@o3c.no